

Lecture 9. Input and output

Functional Programming 2018/19

Alejandro Serrano



Universiteit Utrecht

[Faculty of Science
Information and Computing
Sciences]

Goals

- ▶ Interact with the outside world in Haskell
 - ▶ Input/output and random generation
 - ▶ Learn the difference between pure and impure
- ▶ Introduce `do`-notation

Chapter 10 from Hutton's book



Interactive programs

- ▶ In the olden days, all programs were *batch* programs
 - ▶ Introduce the program and input, sit and drink tea/coffee for hours, and get the output
 - ▶ Programs were isolated from each other
 - ▶ The part of Haskell you have learnt up to now
- ▶ In this modern era, programs are *interactive*
 - ▶ Respond to user input, more like a dialogue
 - ▶ From the perspective of a program, it needs to communicate with an *outside world*
 - ▶ How do we model this in Haskell?



Referential transparency

Referential transparency = you can always substitute a term by its definition without change in the meaning

- ▶ In `let` or `where` bindings:

`let x = e in ... x ... x ...`

is always equivalent to:

`... e ... e ...`

- ▶ In anonymous functions:

`(\x -> e) p`

is always equivalent to:

`let x = p in e`



Referential transparency

A more concrete example:

```
reverse xs ++ xs  
  where xs = filter p ys
```

is equivalent to:

```
reverse (filter p ys) ++ filter p ys
```

Note that the second version duplicates work, but we are speaking here about the *meaning* of the expression, not its efficiency



Referential transparency

- ▶ Referential transparency decouples the meaning of the program from the order of evaluation
 - ▶ Inlining or duplicating does not change the program
- ▶ This has practical advantages:
 - ▶ The compiler can reorder your program for efficiency
 - ▶ Expressions are only evaluated when really needed
 - ▶ This is called *lazy evaluation*
 - ▶ Paralellism becomes much easier



Side-effects

Interaction with the world is not referentially transparent!

Suppose that `getChar :: Char` retrieves the next key stroke from the user

```
let k = getChar in k == k
```

is always `True`, whereas this is not the case with

```
getChar == getChar
```

We say that `getChar` is a **side-effectful** action

- ▶ `getChar` is also called an **impure** function



Side-effects

- ▶ Many other actions have side-effects
 - ▶ Printing to the screen
 - ▶ Generate a random number
 - ▶ Communicate through a network
 - ▶ Talk to a database
- ▶ Intuitively, these actions influence the outside world
 - ▶ Key property: we cannot duplicate the world
 - ▶ And thus we cannot substitute for free



Modelling output

Following this idea, we model an action by a function which changes the world

```
type IO = World -> World
```

Using IO we can give a type to putChar

```
putChar :: Char -> IO  
putChar c world = ... -- details hidden
```



Combining output actions

Executing two actions in sequence is plain composition

```
putAB :: IO
putAB world = putChar 'b' (putChar 'a' world)
-- or using composition
putAB      = putChar 'b' . putChar 'a'
```

The order is not very intuitive



Combining output actions

Executing two actions in sequence is plain composition

```
putAB :: IO
putAB world = putChar 'b' (putChar 'a' world)
-- or using composition
putAB      = putChar 'b' . putChar 'a'
```

The order is not very intuitive

- We introduce *reverse composition*

```
(>>>) :: (a -> b) -> (b -> c) -> a -> c
(f >>> g) x = g (f x)  -- also, flip (.)
```

```
putAB = putChar 'a' >>> putChar 'b'
```



putStr, first version

`putStr s` prints the whole string to the screen

```
putStr :: String -> IO
putStr []      = id  -- keep the world as it is
putStr (c:cs) = putChar c >>> putStr cs
```

`putStrLn s` does the same, with a newline at the end

```
putStrLn s = putStr s >>> putChar '\n'
```



Modelling input

Our IO type is not suitable for `getChar`

- ▶ Solution: pair the output value with the new world

```
type IO a = World -> (a, World)
```

```
getChar :: IO Char
```

```
getChar = ... -- details hidden
```

What is now the return type of `putChar`?

- ▶ We use the empty tuple as a dummy value

```
putChar :: Char -> IO ()
```



Combining input and output

Suppose that we want to echo a character

```
echo = putChar getChar
```

- Couldn't match expected type 'Char'
with actual type 'IO Char'



Combining input and output

Let's try again with function composition

```
echo = getChar >>> putChar
```

- Couldn't match expected type 'IO b'
with actual type 'Char -> IO ()'

```
getChar :: IO Char
```

```
-- World -> (Char, World)
```

```
putChar :: Char -> IO ()
```

```
-- Char -> World -> ((), World)
```

```
(>>>) :: (a -> b) -> (b -> c) -> a -> c
```

Types do not fit, since b should be both (Char, World) –
from getChar – and Char – from putChar



Solution: bind

(>>=) – pronounced “bind” – takes care of threading the world around

`(>>=) :: IO a -> (a -> IO b) -> IO b`

Based on the output of the first action, we choose which action to perform next

```
echo = getChar >>= \c -> putChar c
      -- also getChar >>= putChar
```



Uppercase input

We want to build a `getUpper` function which returns the uppercase version of the last keystroke

```
getChar :: IO Char  
toUpper :: Char -> Char
```

```
getUpper = getChar >>= \c -> toUpper c
```

- Couldn't match expected type 'IO Char' with actual type 'Char'



Uppercase input

We need a way to *embed* pure computations, like `toUpper`, in the impure world

```
return :: a -> IO a
```

Warning! `return` is indeed a very confusing name

- ▶ Does not “break” the flow of the function
- ▶ A more apt synonym is available, `pure`

```
getUpper = getChar >>= \c -> return (toUpper c)  
-- getChar >>= return . toUpper  
-- getChar >>= (toUpper >>> return)
```



Preserving purity

There is no bridge back from the impure to the pure world

```
backFromHell :: IO a -> a
```

In this way we ensure that the outside world never “infects” pure expressions

- ▶ Referential transparency is preserved



Cooking getLine

When dealing with IO, we cannot directly pattern match

- We often use case expressions after (>>=)

```
getLine :: IO String
getLine = getChar >>= (\c ->
    case c of
        '\n' -> return []
        _     -> getLine >>= (\rest ->
            return (c : rest)
        )
    )
```



Cooking getLine

When dealing with IO, we cannot directly pattern match

- We often use case expressions after (>>=)

```
getLine :: IO String
getLine = getChar >>= (\c ->
    case c of
        '\n' -> return []
        -     -> getLine >>= (\rest ->
            return (c : rest)
        )
    )
```

Working directly with (>>=) is very cumbersome!



do-notation

Luckily, Haskell has specific notation for IO

```
getLine = do c <- getChar
           case c of
             '\n' -> return []
             _     -> do rest <- getLine
                       return (c : rest)
```

Blocks for IO start with the keyword `do`

- ▶ `<-` gives a name to the result of an IO action
- ▶ The notation was chosen to “look imperative”



Cooking putStr

Let us rewrite putStr with the new combinators

```
putStr :: String -> IO ()  
putStr []      = return ()  
putStr (c:cs) = putChar c >>= (\_ -> putStr cs)
```

What is happening is much clearer with do-notation

```
putStr :: String -> IO ()  
putStr []      = return ()  
putStr (c:cs) = do putChar c  
                  putStr cs
```



do-notation, in general

A general do block is translated as nested ($\gg=$)

<code>do</code>	<code>x1 <- a1</code>		<code>a1 >>= (\x1 -></code>
	<code>x2 <- a2</code>		<code>a2 >>= (\x2 -></code>
	<code>...</code>	<code>==></code>	<code>...</code>
	<code>xn <- an</code>		<code>an >>= (\xn -></code>
	<code>expr</code>		<code>expr) ...))</code>

In addition, if you don't care about a value, you can write simply `ai` instead of `_ <- ai`

Rule of thumb: do not think about ($\gg=$) at all, just use `do`



Guess a number

Pick a number between 1 and 100.

Is it 50? (g = greater, l = less, c = correct)

g

Is it 75? (g = greater, l = less, c = correct)

l

Is it 62? (g = greater, l = less, c = correct)

g

Is it 68? (g = greater, l = less, c = correct)

l

Is it 65? (g = greater, l = less, c = correct)

c

Guessed



Guess a number

We do *binary search* over the list of numbers

- At each step, we pick the middle value as a guess

```
guess :: Int -> Int -> IO ()
guess l u
  = do let m = (u + l) `div` 2
        putStr ("Is it " ++ show m ++ "?")
        putStrLn "(g = greater, l = less, c = correct)"
        k <- getChar
        case k of
          'g' -> guess (m + 1) u
          'l' -> guess l (m - 1)
          'c' -> putStrLn "Guessed"
          _   -> do putStrLn "Press type g/l/c!"
                  guess l u
```



Guess a number, main program

When an executable written in Haskell starts, the `main` function is called

- ▶ `main` always has type `IO ()`

```
main :: IO ()  
main = do (l:u:_) <- getArgs  
          guess (read l) (read u)
```

- ▶ `getArgs :: IO [String]` obtains program arguments
- ▶ `read :: Read a => String -> a`
 - ▶ Parses a `String` into a value
 - ▶ In this case, we parse it into an `Int`



Summary of basic I/O actions

```
return    :: a -> IO a
```

```
getChar   :: IO Char
```

```
getLine   :: IO String
```

```
putChar   :: Char -> IO ()
```

```
putStr    :: String -> IO ()
```

```
putStrLn  :: String -> IO ()
```



Dealing with files

The simplest functions to work with files in Haskell

```
type FilePath = String
```

```
readFile  :: FilePath -> IO String
```

```
writeFile :: FilePath -> String -> IO ()
```

The following functions are often convenient

```
lines    :: String -> [String]    -- break at '\n'
```

```
unlines  :: [String] -> String    -- join lines
```

```
-- convert back and forth
```

```
show :: Show a => a -> String
```

```
read  :: Read a => String -> a
```



Guess a number, bounds from file

```
main :: IO ()
main = do -- Read the contents of the file
          config <- readFile "guess.config"
          -- Get the first two lines
          let l:u:_ = lines config
          -- Parse the numbers and start guessing
          guess (read l) (read u)
```



IO as first-class citizens



Universiteit Utrecht

[Faculty of Science
Information and Computing
Sciences]

IO actions are first-class

In the same way as you do with functions

- ▶ An IO action can be an argument or result of a function
- ▶ IO actions can be put in a list or other container

```
map (\name -> putStrLn ("Hello, " ++ name))  
  ["Mary", "John"]  :: [IO ()]
```



Building versus execution of IO actions

```
map (\name -> putStrLn ("Hello, " ++ name))  
    ["Mary", "John"]  ::  [IO ()]
```

Running this code prints **nothing** to the screen

- ▶ We say that it *builds* the IO actions: describes what needs to be done but does not do it yet

To obtain the side-effects, you need to *execute* the actions

- ▶ At the interpreter prompt
- ▶ In a `do` block which is ultimately called by `main`
- ▶ An executed action always has a `IO T` type



Execute a list of actions

`sequence_` as performs the side-effects of a list of actions

1. Define the type

```
sequence_ :: [IO a] -> IO ()
```

2. Enumerate the cases

```
sequence_ []          = _  
sequence_ (a:as)     = _
```

3. Define the cases

```
sequence_ []          = return ()  
sequence_ (a:as)     = do a  
                      sequence_ as
```



Execute a list of actions

We have all the ingredients to greet a list of people

```
greet :: [String] -> IO ()  
greet = sequence_  
    . map (\name -> putStrLn ("Hello, " ++ name))
```

This combination is very common, so the library defines

```
mapM_ :: (a -> IO b) -> [a] -> IO ()  
  
greet = mapM_ (\name -> putStrLn ("Hello, " ++ name))
```



Execute a list of actions

By just flipping the order of arguments, we can write “imperative-looking” code

```
forM_ :: [a] -> (a -> IO b) -> IO ()  
forM_ = flip mapM_
```

```
greet names = forM_ names $ \name ->  
                putStrLn ("Hello, " ++ name)
```



Answer to a yes-no questions

`poseQuestion q` prints a question to the screen, obtains a `y` or `n` input from the user and returns it as a Boolean

```
poseQuestion :: String -> IO Bool
poseQuestion q
  = do putStr q
       putStrLn "Answer (y) or (n)"
       (k:_) <- getLine
       case k of
         'y' -> return True
         'n' -> return False
         _   -> do putStrLn "Answer (y) or (n)"
                   poseQuestion q
```



Gathering all answers

Once again, if we map over the list the actions are inside

```
map poseQuestion qs :: [IO Bool]
```

sequence_ does not work, since it throws away the result

```
sequence :: [IO a] -> IO [a]
sequence []      = return []
sequence (a:as) = do r  <- a
                    rs <- sequence as
                    return (r:rs)
```



Gathering all answers

Now we can gather answers to all questions at once

```
poseQuestions :: [String] -> IO [Bool]
poseQuestions = sequence . map poseQuestion
```

We have non-forgetful versions of the previous functions

```
mapM :: (a -> IO b) -> [a] -> IO [b]
forM :: [a] -> (a -> IO b) -> IO [b]
```

Naming convention: a function which ends in `_` throws away information



Lifting, the problem

Take the last equation of sequence

```
sequence (a:as) = do r  <- a  
                  rs <- sequence as  
                  return (r:rs)
```

Or a function which reads two numbers and returns its sum

```
inputSum = do x <- getLine  
             y <- getLine  
             return (read x + read y)
```

In both cases we are just gathering a sequence of values from actions and applying a (pure) function at the very end



Lifting, the solution

From a previous lecture

How do we make (+) work with Maybe values?

```
liftM2 :: (a -> b -> c)
        -> Maybe a -> Maybe b -> Maybe c
```



Lifting, the solution

From a previous lecture

How do we make (+) work with Maybe values?

```
liftM2 :: (a -> b -> c)
        -> Maybe a -> Maybe b -> Maybe c
```

The same function is available for IO!

```
liftM2 :: (a -> b -> c)
        -> IO a -> IO b -> IO c
```



Lifting, examples

The code becomes much shorter!

```
sequence (a:as) = liftM2 (:) a (sequence as)
```



Lifting, examples

The code becomes much shorter!

```
sequence (a:as) = liftM2 (:) a (sequence as)
```

In fact, since IO is a monad, people often write

```
sequence (a:as) = (:) <$> a <*> sequence as
```

This style of programming is covered later in the course



Randomness



Universiteit Utrecht

[Faculty of Science
Information and Computing
Sciences]

Random generation

Random generation is provided by the `System.Random` module of the `random` package

```
class Random a where
  randomR :: RandomGen g => (a, a) -> g -> (a, g)
  random  :: RandomGen g =>                g -> (a, g)
```

- ▶ `a` is the type of value you want to generate
- ▶ `g` is the type of random generators
 - ▶ Usually, random generators keep some additional information called the *seed*



Generating several random numbers

If you want to generate several values, you need to keep track of the seed yourself

```
generateTwoNumbers :: RandomGen g
                  => g -> ((Int, Int), g)

generateTwoNumbers g
  = let (v1, g1) = random g
        (v2, g2) = random g1  -- Use new seed
    in ((v1, v2), g2)         -- Return last seed
```



Obtaining the seed

An initial value for the generator needs external input

- ▶ The following function takes care of obtaining a new seed, performing random generation and updating the seed at the end

```
getStdRandom :: (StdGen -> (a, StdGen)) -> IO a
```

- ▶ Note the use of a higher-order function to encapsulate the part of the program which needs randomness

Because of its ubiquity, the following functions are provided

```
randomRIO = getStdRandom . randomR
```

```
randomIO  = getStdRandom    random
```



Summary

- ▶ Actions with side-effects which return a value of type `a` are represented by `IO a`
 - ▶ Pure and impure parts are perfectly delineated
 - ▶ The `main` in a Haskell program has type `IO ()`
- ▶ To sequence IO actions, use `do`-notation
 - ▶ Under the hood it translates to nested `(>>=)` (bind)
- ▶ IO actions are first-class citizens

