



Lazy evaluation

Functional Programming

Haskell can be defined with four adjectives

- Functional
- Statically typed
- Pure
- Lazy

Haskell can be defined with four adjectives

- Functional
- Statically typed
- Pure
- **Lazy**

Goals

- Understand the lazy evaluation strategy
 - As opposed to strict evaluation
- Understand why laziness is useful
 - ...
 - Work with infinite structures
- Learn about laziness pitfalls
 - Force evaluation using seq

A simple expression

```
square :: Integer -> Integer
```

```
square x = x * x
```

```
square (1 + 2)
```

```
= -- magic happens in the computer
```

```
9
```

How do we reach that final value?

Strict or eager or call-by-value evaluation

In most programming languages:

1. Evaluate the arguments completely
2. Evaluate the function call

```
square (1 + 2)
```

```
= -- evaluate arguments
```

```
square 3
```

```
= -- go into the function body
```

```
3 * 3
```

```
=
```

```
9
```

Non-strict or call-by-name evaluation

Arguments are replaced as-is in the function body

```
square (1 + 2)
```

```
= -- go into the function body
```

```
(1 + 2) * (1 + 2)
```

```
= -- we need the value of (1 + 2) to continue
```

```
3 * (1 + 2)
```

```
=
```

```
3 * 3
```

```
=
```

```
9
```

Does call-by-name make any sense?

In the case of square, non-strict evaluation is worse

Is this always the case?

Does call-by-name make any sense?

In the case of square, non-strict evaluation is worse

Is this always the case?

```
const x y = x  -- forget about y
```

```
-- Call-by-value
```

```
const 5 (1 + 2)
```

```
=
```

```
const 5 3
```

```
=
```

```
5
```

```
-- Call-by-name
```

```
const 5 (1 + 2)
```

```
=
```

```
5
```

Sharing expressions

square (1 + 2)

=

(1 + 2) * (1 + 2)

Why redo the work for (1 + 2)?

Sharing expressions

square (1 + 2)

=

(1 + 2) * (1 + 2)

Why redo the work for (1 + 2)?

We can share the evaluated result

square (1 + 2)

=

Δ * Δ

\uparrow \uparrow (1 + 2)

= 3

=

9

Lazy evaluation

Haskell uses a **lazy** evaluation strategy

- Expressions are not evaluated *until needed*
- Duplicate expressions are *shared*

Lazy evaluation never requires more steps than call-by-value

Each of those not-evaluated expressions is called a **thunk**

Does it matter?

Is it possible to get different outcomes using different evaluation strategies?

Does it matter?

Is it possible to get different outcomes using different evaluation strategies?

No and Yes

Does it matter?

- No:

Theorem [Church-Rosser Theorem]

For *terminating* programs all evaluation strategies produce the same result value.

Does it matter?

- No:

Theorem [Church-Rosser Theorem]

For *terminating* programs all evaluation strategies produce the same result value.

- Yes:

Does it matter?

- No:

Theorem [Church-Rosser Theorem]

For *terminating* programs all evaluation strategies produce the same result value.

- Yes:

1. Holds only for terminating programs.

- What about infinite loops?
- What about exceptions?

Does it matter?

- No:

Theorem [Church-Rosser Theorem]

For *terminating* programs all evaluation strategies produce the same result value.

- Yes:

1. Holds only for terminating programs.

- What about infinite loops?
- What about exceptions?

2. Performance might be different.

- As square and const show

Termination

`loop x = loop x`

- This is a well-typed program
- But `loop 3` never terminates

Termination

`loop x = loop x`

- This is a well-typed program
- But `loop 3` never terminates

Question: What does `'const 5 (loop 3)'` evaluate to?

Termination

`loop x = loop x`

- This is a well-typed program
- But `loop 3` never terminates

Question: What does '`const 5 (loop 3)`' evaluate to?

`-- Eager`

`const 5 (loop 3)`

`=`

`const 5 (loop 3)`

`=`

`...`

`-- Lazy`

`const 5 (loop 3)`

`=`

`5`

Observation:

Lazy evaluation terminates more often than eager evaluation.

Question: Why is this useful?

Short-circuiting

```
(&&)      :: Bool -> Bool -> Bool
```

```
False && _ = False
```

```
True  && x = x
```

- In eager languages, `x && y` evaluates both conditions
 - But if the first one fails, why bother?
 - C/Java/C# include a built-in *short-circuit* conjunction
- In Haskell, `x && y` only evaluates the second argument if the first one is `True`
 - `False && (loop True)` terminates

Why? Build your own Control structures

```
if_      :: Bool -> a -> a -> a
if_ True  t _ = t
if_ False _ e = e
```

- In eager languages, `if_` evaluates both branches
- In lazy languages, only the one being selected

Why? Build your own Control structures

```
if_      :: Bool -> a -> a -> a
if_ True  t _ = t
if_ False _ e = e
```

- In eager languages, `if_` evaluates both branches
- In lazy languages, only the one being selected

For that reason,

- In eager languages, `if` has to be *built-in*
- In lazy languages, you can build your *own control structures*

Why? Separation of Concerns

- Lazyness allows for easier separation of concerns.

```
data Operation = Sum | Product
```

```
apply      :: Operation -> [Int] -> Int
```

```
apply op xs = case op of  
    Sum      -> sumResult  
    Product  -> productResult
```

```
where
```

```
    sumResult      = sum xs
```

```
    productResult = product xs
```

Why? Separation of Concerns

- Lazyness allows for easier separation of concerns.

```
minAndMax :: Ord a => a -> [a] -> (a,a)
```

```
minimum'   :: Ord a => a -> [a] -> a
```

```
minimum' d = fst . minAndMax d
```

Why? Infinite structures

An infinite list of ones:

```
ones :: [Integer]
```

```
ones = 1 : ones
```

ones is infinite, but everything works fine if we only work with a *finite* part

```
take 2 ones  
= take 2 (1 : ones)  
= 1 : take 1 ones  
= 1 : take 1 (1 : ones)  
= 1 : 1 : take 0 ones  
= 1 : 1 : []
```

A list of all natural numbers

To build an infinite list of numbers, we use recursion

- This kind of recursion is trickier than the usual one

```
nats :: [Integer]
```

```
nats = 0 : map (+1) nats
```

```
    take 2 nats
```

```
= take 2 (0 : map (+1) nats)
```

```
= 0 : take 1 (map (+1) nats)
```

```
= 0 : take 1 (map (+1) (0 : map (+1) nats))
```

```
= 0 : take 1 (1 : map (+1) (map (+1) nats))
```

```
= 0 : 1 : take 0 (map (+1) (map (+1) nats))
```

```
= 0 : 1 : []
```

A list of all Fibonacci numbers

Remember the usual definition of fib,

`fib 0 = 0`

`fib 1 = 1`

`fib n = fib (n-1) + fib (n-2)`

A list of all Fibonacci numbers

Remember the usual definition of fib,

```
fib 0 = 0
```

```
fib 1 = 1
```

```
fib n = fib (n-1) + fib (n-2)
```

Here is a list containing all Fibonacci numbers:

```
fibs :: [Integer]
```

```
fibs = 0 : 1 : zipWith (+) fibs (tail fibs)
```

```
fib :: Integer -> Integer
```

```
fib n = fibs !! n  -- Take the n-th element
```

A list of all Fibonacci numbers

0 : 1 : ...

+ 1 : ...

1 : ...

A list of all Fibonacci numbers

$$\begin{array}{rccccccc} & 0 & : & 1 & : & 1 & : & \dots \\ + & 1 & : & 1 & : & & & \dots \\ \hline & 1 & : & 2 & : & & & \dots \end{array}$$

A list of all Fibonacci numbers

	0	:	1	:	1	:	2	:	...
+	1	:	1	:	2	:	...		

	1	:	2	:	3	:	...		

A list of all prime numbers: Sieve of Erastosthenes

An algorithm to compute the list of all primes

- Already known in Ancient Greece
1. Lay all numbers in a list starting with 2
 2. Take the first next number p in the list
 3. Remove all the multiples of p from the list
 - $2p, 3p, 4p...$
 - Alternatively, remove n if the remainder with p is 0
 4. Go back to step 2 with the first remaining number

Sieve of Erastosthenes

1. Lay all numbers in a list starting with 2

```
primes :: [Integer]
primes = sieve [2 .. ]  -- an infinite list
```

Sieve of Erastosthenes

1. Lay all numbers in a list starting with 2

```
primes :: [Integer]
primes = sieve [2 .. ]  -- an infinite list
```

2. Take the first number p in the list

```
sieve (p:ns) = p : ...
```

3. Remove n if the remainder with p is 0

4. Go back to step 2 with the first remaining number

```
sieve (p:ns)
  = p : sieve [n | n <- ns, n `mod` p /= 0]
```

“Until needed”

How does Haskell know *how much* to evaluate?

- By default, everything is kept in a thunk
- When we have a case distinction, we evaluate enough to distinguish which branch to follow

```
take 0 _      = []
```

```
take _ []     = []
```

```
take n (x:xs) = x : take (n-1) xs
```

- If the number is 0 we do not need the list at all
- Otherwise, we need to distinguish [] from x:xs

Weak Head Normal Form

An expression is in **weak head normal form** (WHNF) if it is:

- A constructor with (possibly non-evaluated) data inside
 - `True` or `Just (1 + 2)`
- An anonymous function
 - The body might be in any form
 - `\x -> x + 1` or `\x -> if_ True x x`
- A function applied to too few arguments
 - `map minimum`

Every time we need to distinguish the branch to follow the expression is evaluated until its WHNF

Weak Head Normal Form

Which of these expressions are in WHNF?

1. `zip [1..]`
2. `Node Leaf 4 (fmap (+1) Leaf)`
3. `map (x:) xs`
4. `height (Node Leaf 'a' (Node Leaf 'b' Leaf))`
5. `_ b -> b`
6. `map (\x -> x + 1) [1..5]`
7. `(x + 1) : foldr (:) [] [1..5]`

Weak Head Normal Form

Which of these expressions are in WHNF?

1. `zip [1..]`
2. `Node Leaf 4 (fmap (+1) Leaf)`
3. `map (x:) xs`
4. `height (Node Leaf 'a' (Node Leaf 'b' Leaf))`
5. `_ b -> b`
6. `map (\x -> x + 1) [1..5]`
7. `(x + 1) : foldr (:) [] [1..5]`

answer: 1,2,5,7

Strict versus lazy functions

Note the difference between these two functions

```
loop 2 + 3
= -- definition of loop
  loop 2 + 3
= -- never-ending sequence
  ...
```

```
const 3 (loop 2)
= -- definition of const
  3
  -- and that's it!
```

Strict versus lazy functions

A function is **strict** on one argument if the result of the function is non-terminating given a non-terminating value for that argument

- $(+)$ is strict on its first and second arguments
- `const` is not strict on its second argument, but strict on the first

We represent non-termination by \perp or undefined

- We also call \perp a *diverging* computation
- f is strict if $f \perp = \perp$

Some (tricky) questions

What is the result of these expressions?

1. `(\x -> x) True`
2. `(\x -> x) undefined`
3. `(\x -> 0) undefined`
4. `(\x -> undefined) 0`
5. `(\x f -> f x) undefined`
6. `undefined undefined`
7. `length (map undefined [1,2])`

Some (tricky) questions

What is the result of these expressions?

1. `(\x -> x) True` = `True`
2. `(\x -> x) undefined` = `undefined`
3. `(\x -> 0) undefined` = `0`
4. `(\x -> undefined) 0` = `undefined`
5. `(\x f -> f x) undefined` = `\f -> f undefined`
6. `undefined undefined` = `undefined`
7. `length (map undefined [1,2])` = `2`

Lazy Evaluation vs Performance

Case study: foldl

From a long, long time ago...

```
foldl _ v []      = v
```

```
foldl f v (x:xs) = foldl f (f v x) xs
```

Case study: foldl

From a long, long time ago...

```
foldl _ v [] = v
```

```
foldl f v (x:xs) = foldl f (f v x) xs
```

```
foldl (+) 0 [1,2,3]
```


Case study: foldl

From a long, long time ago...

```
foldl _ v [] = v
```

```
foldl f v (x:xs) = foldl f (f v x) xs
```

```
foldl (+) 0 [1,2,3]
```

```
= foldl (+) (0 + 1) [2,3]
```

```
= foldl (+) ((0 + 1) + 2) [3]
```

```
= foldl (+) (((0 + 1) + 2) + 3) []
```

```
= ((0 + 1) + 2) + 3
```

Case study: foldl

```
foldl (+) 0 [1,2,3]  
= ((0 + 1) + 2) + 3
```

Question: What is the problem with this?


Case study: foldl

```
foldl (+) 0 [1,2,3]  
= ((0 + 1) + 2) + 3
```

Question: What is the problem with this?

- Each of the additions is kept in a thunk
 - Some memory need to be reserved!

Case study: foldl



Space leak = data structure which grows bigger, or lives longer than expected

- More memory in use means more *Garbage Collection*
- As a result, performance decreases

The most common source of space leaks are thunks

- Thunks are essential for lazy evaluation
- But they also take some amount of memory

Garbage collection

- Thunks are managed by the run-time system
 - They are created when you need a value
 - But are not reclaimed right after evaluation
- Haskell uses **garbage collection** (GC)
 - Every now and then Haskell takes back all the memory used by thunks which are not needed anymore
 - *Pro*: we do not need to care about memory
 - *Con*: GC takes time, so lags can occur
- Most modern languages nowadays use GC
 - Java, Scala, C#, Ruby, Python...
 - Swift uses Automatic Reference Counting (ARC)

Case study: foldl

We want to reduce memory usage and speed up the computation.

We *force* additions before going on

```
foldl (+) 0 [1,2,3]
= foldl (+) (0 + 1) [2,3]
= foldl (+) 1 [2,3]
= foldl (+) (1 + 2) [3]
= foldl (+) 3 [3]
= foldl (+) (3 + 3) []
= foldl (+) 6 []
= 6
```

Forcing evaluation

Haskell has a primitive operation to force

```
seq :: a -> b -> b
```

A call of the form `seq x y`

- First evaluates `x` up to WHNF
- Then it proceeds normally to compute `y`

Usually, `y` depends on `x` somehow


Case study: foldl

We can write a new version of `foldl` which forces the accumulated value before recursion is unfolded

```
foldl' _ v []      = v
foldl' f v (x:xs) = let z = f v x
                   in z `seq` foldl' f z xs
```

This version solves the problem with addition

Case study: foldl



Strict application

Most of the times we use `seq` to force an argument to a function, that is, *strict application*

```
($!) :: (a -> b) -> a -> b
```

```
f $! x = x `seq` f x
```

Because of sharing, `x` is evaluated only once

More (tricky) questions

What is the result of these expressions?

1. `(\x -> 0) $! undefined`
2. `seq (undefined, undefined) 0`
3. `snd $! (undefined, undefined)`
4. `(\x -> 0) $! (\x -> undefined)`
5. `undefined $! undefined`
6. `length $! map undefined [1,2]`
7. `seq (undefined + undefined) 0`
8. `seq (foldr undefined undefined) 0`
9. `seq (1 : undefined) 0`

More (tricky) questions

What is the result of these expressions?

1. `(\x -> 0) $! undefined = undefined`
2. `seq (undefined, undefined) 0 = 0`
3. `snd $! (undefined, undefined) = undefined`
4. `(\x -> 0) $! (\x -> undefined) = 0`
5. `undefined $! undefined = undefined`
6. `length $! map undefined [1,2] = 2`
7. `seq (undefined + undefined) 0 = undefined`
8. `seq (foldr undefined undefined) 0 = 0`
9. `seq (1 : undefined) 0 = 0`

seq only evaluates up to WHNF

Case study: Fibonacci numbers


`fib 0 = 0`

`fib 1 = 1`


`fib n = fib (n-1) + fib (n-2)`

What happens when we ask for `fib 5`?

Case study: Fibonacci numbers



Case study: Fibonacci numbers



Local memoization (aka Dynamic Programming)

Idea: remember the result for function calls

- We build a list of partial results
- Sharing takes care of evaluating only once

```
memo_fib n = go n
  where go i  = fibs !! i
        fibs  = map fib [0 .. ]
        fib 0 = 0
        fib 1 = 1
        fib n = go (n-1) + go (n-2)
```

You can get even faster by using a better data structure

- For example, IntMap from containers

Summary

- Laziness = evaluate only as much as needed
 - As opposed to the more common *eager* evaluation
- Evaluation is guided by pattern matching
 - We need WHNF to choose a branch
 - Some arguments may not even be evaluated
- Laziness is tricky when it fails
 - Too many thunks lead to a space leak
 - `seq` is used to *force* evaluation