

Lecture 11. Lazy evaluation

Functional Programming 2018/19

Alejandro Serrano



Universiteit Utrecht

[Faculty of Science
Information and Computing
Sciences]

Goals

- ▶ Understand the lazy evaluation strategy
 - ▶ As opposed to strict evaluation
- ▶ Work with infinite structures
- ▶ Learn about laziness pitfalls
 - ▶ Force evaluation using `seq`



A simple expression

```
square :: Integer -> Integer
```

```
square x = x * x
```

```
square (1 + 2)
```

```
= -- magic happens in the computer
```

```
9
```

How do we reach that final value?



Strict or eager or call-by-value evaluation

In most programming languages:

1. Evaluate the arguments completely
2. Evaluate the function call

```
square (1 + 2)
= -- evaluate arguments
square 3
= -- go into the function body
3 * 3
=
9
```



Non-strict or call-by-name evaluation

Arguments are replaced as-is in the function body

```
square (1 + 2)
= -- go into the function body
  (1 + 2) * (1 + 2)
= -- we need the value of (1 + 2) to continue
  3 * (1 + 2)
=
  3 * 3
=
  9
```



Does call-by-name make any sense?

In the case of `square`, non-strict evaluation is worse

Is this always the case?



Does call-by-name make any sense?

In the case of `square`, non-strict evaluation is worse

Is this always the case?

```
const x y = y  -- forget about x
```

```
-- Call-by-value
```

```
const (1 + 2) 5
```

```
=
```

```
const 3 5
```

```
=
```

```
5
```

```
-- Call-by-name
```

```
const (1 + 2) 5
```

```
=
```

```
5
```



Sharing expressions

square (1 + 2)

=

(1 + 2) * (1 + 2)

Why redo the work for (1 + 2)?



Sharing expressions

```
square (1 + 2)
=
(1 + 2) * (1 + 2)
```

Why redo the work for (1 + 2)?

We can share the evaluated result

```
square (1 + 2)
=
Δ * Δ
↑___↑___ (1 + 2)
      = 3
=
9
```



Lazy evaluation

Haskell uses a **lazy** evaluation strategy

- ▶ Expressions are not evaluated *until needed*
- ▶ Duplicate expressions are *shared*

Lazy evaluation never requires more steps than call-by-value

Each of those not-evaluated expressions is called a **thunk**



Does it matter?

Is it possible to get different outcomes using different evaluation strategies?

...

Yes and no



Does it matter? - Correctness and efficiency

The *Church-Rosser Theorem* states that for *terminating* programs the result of the computation does *not* depend on the evaluation strategy

But...

1. Performance might be different
 - ▶ As square and const show
2. This applies only if the program terminates
 - ▶ What about infinite loops?
 - ▶ What about exceptions?



Termination

```
loop x = loop x
```

- ▶ This is a well-typed program
- ▶ But `loop 3` never terminates

```
-- Eager
```

```
const (loop 3) 5
```

```
=
```

```
const (loop 3) 5
```

```
=
```

```
...
```

```
-- Lazy
```

```
const (loop 3) 5
```

```
=
```

```
5
```

Lazy evaluation terminates more often than eager



Build your own control structures

```
if_ :: Bool -> a -> a -> a
if_ True  t _ = t
if_ False _ e = e
```

- ▶ In eager languages, `if_` evaluates both branches
- ▶ In lazy languages, only the one being selected

For that reason,

- ▶ In eager languages, `if` has to be *built-in*
- ▶ In lazy languages, you can build your *own control structures*



Short-circuiting

```
( $\&\&$ ) :: Bool -> Bool -> Bool
```

```
False  $\&\&$  _ = False
```

```
True  $\&\&$  x = x
```

- ▶ In eager languages, $x \ \&\& \ y$ evaluates both conditions
 - ▶ But if the first one fails, why bother?
 - ▶ C/Java/C# include a built-in *short-circuit* conjunction
- ▶ In Haskell, $x \ \&\& \ y$ only evaluates the second argument if the first one is `True`
 - ▶ `False && (loop True)` terminates



An infinite list of ones

```
ones :: [Integer]
ones = 1 : ones
```

`ones` is infinite, but everything works fine if we only work with a *finite* part

```
take 2 ones
= take 2 (1 : ones)
= 1 : take 1 ones
= 1 : take 1 (1 : ones)
= 1 : 1 : take 0 ones
= 1 : 1 : []
```



A list of all natural numbers

To build an infinite list of numbers, we use recursion

- This kind of recursion is trickier than the usual one

```
nats :: [Integer]
nats = 0 : map (+1) nats
```

```
take 2 nats
= take 2 (0 : map (+1) nats)
= 0 : take 1 (map (+1) nats)
= 0 : take 1 (map (+1) (0 : map (+1) nats))
= 0 : take 1 (1 : map (+1) (map (+1) nats))
= 0 : 1 : take 0 (map (+1) (map (+1) nats))
= 0 : 1 : []
```



“Until needed”

How does Haskell know *how much* to evaluate?

- ▶ By default, everything is kept in a thunk
- ▶ When we have a case distinction, we evaluate enough to distinguish which branch to follow

```
take 0 _      = []
```

```
take _ []     = []
```

```
take n (x:xs) = x : take (n-1) xs
```

- ▶ If the number is 0 we do not need the list at all
- ▶ Otherwise, we need to distinguish [] from x:xs



Weak Head Normal Form

An expression is in **weak head normal form** (WHNF) if it is:

- ▶ A constructor with (possibly non-evaluated) data inside
 - ▶ `True` or `Just (1 + 2)`
- ▶ An anonymous function
 - ▶ The body might be in any form
 - ▶ `\x -> x + 1` or `\x -> if_ True x x`
- ▶ A built-in function applied to too few arguments

Every time we need to distinguish the branch to follow the expression is evaluated until its WHNF



A list of all Fibonacci numbers

Remember the usual definition of `fib`,

```
fib 0 = 1
```

```
fib 1 = 1
```

```
fib n = fib (n-1) + fib (n-2)
```



A list of all Fibonacci numbers

Remember the usual definition of `fib`,

```
fib 0 = 1
```

```
fib 1 = 1
```

```
fib n = fib (n-1) + fib (n-2)
```

Here is a list containing all Fibonacci numbers:

```
fibs :: [Integer]
```

```
fibs = 0 : 1 : zipWith (+) fibs (tail fibs)
```

```
fib :: Integer -> Integer
```

```
fib n = fibs !! n  -- Take the n-th element
```



A list of all Fibonacci numbers

$$\begin{array}{rcll} & 0 & : & 1 & : & \dots \\ + & 1 & : & \dots & & \\ \hline & 1 & : & \dots & & \end{array}$$



A list of all Fibonacci numbers

$$\begin{array}{rccccccc} & 0 & : & 1 & : & 1 & : & \dots \\ + & 1 & : & 1 & : & & & \dots \\ \hline & 1 & : & 2 & : & & & \dots \end{array}$$



A list of all Fibonacci numbers

$$\begin{array}{rccccccccc} & 0 & : & 1 & : & 1 & : & 2 & : & \dots \\ + & 1 & : & 1 & : & 2 & : & \dots & & \\ \hline & 1 & : & 2 & : & 3 & : & \dots & & \end{array}$$



Sieve of Erastosthenes

An algorithm to compute the list of all primes

- ▶ Already known in Ancient Greece

1. Lay all numbers in a list starting with 2
2. Take the first next number p in the list
3. Remove all the multiples of p from the list
 - ▶ $2p, 3p, 4p \dots$
 - ▶ Alternatively, remove n if the remainder with p is 0
4. Go back to step 2 with the first remaining number



Sieve of Erastosthenes

1. Lay all numbers in a list starting with 2

```
primes :: [Int]
primes = sieve [2 .. ]  -- an infinite list
```

2. Take the first number p in the list

```
sieve (p:ns) = ...
```

3. Remove n if the remainder with p is 0

4. Go back to step 2 with the first remaining number

```
sieve (p:ns)
  = p : sieve [n | n <- ns, n `mod` p /= 0]
```



Strict versus lazy functions

Note the difference between these two functions

```
loop 2 + 3
= -- definition of loop
loop 2 + 3
= -- never-ending sequence
...
```

```
const 3 (loop 2)
= -- definition of const
3
-- and that's it!
```



Strict versus lazy functions

A function is **strict** on one argument if the result of the function is non-terminating given a non-terminating value for that argument

- ▶ `(+)` is strict on its first and second arguments
- ▶ `const` is not strict on its second argument, but strict on the first

We represent non-termination by \perp or `undefined`

- ▶ We also call \perp a *diverging* computation
- ▶ f is strict if $f \perp = \perp$



Some (tricky) questions

What is the result of these expressions?

1. `(\x -> x) True`
2. `(\x -> x) undefined`
3. `(\x -> 0) undefined`
4. `(\x -> undefined) 0`
5. `(\x f -> f x) undefined`
6. `undefined undefined`
7. `length (map undefined [1,2])`



Some (tricky) questions

What is the result of these expressions?

1. `(\x -> x) True` = `True`
2. `(\x -> x) undefined` = `undefined`
3. `(\x -> 0) undefined` = `0`
4. `(\x -> undefined) 0` = `undefined`
5. `(\x f -> f x) undefined` = `\f -> f undefined`
6. `undefined undefined` = `undefined`
7. `length (map undefined [1,2])` = `2`



Garbage collection

- ▶ Thunks are managed by the run-time system
 - ▶ They are created when you need a value
 - ▶ But are not reclaimed right after evaluation
- ▶ Haskell uses **garbage collection** (GC)
 - ▶ Every now and then Haskell takes back all the memory used by thunks which are not needed anymore
 - ▶ *Pro*: we do not need to care about memory
 - ▶ *Pro*: GC enables fancy distributed algorithms
 - ▶ *Con*: GC takes time, so lags can occur
- ▶ Most modern languages nowadays use GC
 - ▶ Java, Scala, C#, Ruby, Python...
 - ▶ Swift uses Automatic Reference Counting (ARC)



Performance



Universiteit Utrecht

[Faculty of Science
Information and Computing
Sciences]

Lists are sloooooooooow

You have to traverse them for almost every operation

- ▶ To find an element you need to compare all the ones in front of it
- ▶ To append two lists you traverse the first one

We have seen some techniques that help

- ▶ Use an accumulator for reversing a list
- ▶ Use a search tree to find things quickly

Still, lists are convenient for small data sets



Use better data structures

`containers` contains many general purpose data structures

- ▶ `Map k v` hold values `v` indexed by keys of type `k`
- ▶ `Set a` holds values of `a` without repetition
- ▶ `Seq a` is similar to a list, but more efficient
 - ▶ Match and build from both sides
 - ▶ Useful as a queue or as a stack
- ▶ `Tree a` implement rose trees

The interface is almost *identical* to that of lists



Use better data structures

String is just a synonym for `[Char]`

- ▶ Simple to handle, very inefficient

If your application uses a lot of them, you should consider

- ▶ `ByteString` to treat it as an array of words
- ▶ `Text` to represent Unicode strings



Space leaks

Space leak = data structure which grows bigger, or lives longer than expected

- ▶ More memory in use means more GC
- ▶ As a result, performance decreases

The most common source of space leaks are thunks

- ▶ Thunks are essential for lazy evaluation
- ▶ But they also take some amount of memory



Case study: foldl'

From long, long time ago...

```
foldl _ v []      = v
foldl f v (x:xs) = foldl f (f v x) xs
```

```
foldl (+) 0 [1,2,3]
= foldl (+) (0 + 1) [2,3]
= foldl (+) ((0 + 1) + 2) [3]
= foldl (+) (((0 + 1) + 2) + 3) []
= ((0 + 1) + 2) + 3
```



Case study: foldl'

```
foldl (+) 0 [1,2,3]  
= ((0 + 1) + 2) + 3
```

- ▶ Each of the additions is kept in a thunk
 - ▶ Some memory need to be reserved
 - ▶ They have to be GC'ed after use



Case study: foldl'




Figure 1:

[Faculty of Science
Information and Computing
Sciences]



Universiteit Utrecht

Case study: foldl'

Just performing the addition is faster!

- ▶ Computers are fast at arithmetic
- ▶ We want to *force* additions before going on

```
foldl (+) 0 [1,2,3]
= foldl (+) (0 + 1) [2,3]
= foldl (+) 1 [2,3]
= foldl (+) (1 + 2) [3]
= foldl (+) 3 [3]
= foldl (+) (3 + 3) []
= foldl (+) 6 []
= 6
```



Forcing evaluation

Haskell has a primitive operation to force

`seq :: a -> b -> b`

A call of the form `seq x y`

- ▶ First evaluates `x` up to WHNF
- ▶ Then it proceeds normally to compute `y`

Usually, `y` depends on `x` somehow



Case study: foldl'

We can write a new version of `foldl` which forces the accumulated value before recursion is unfolded

```
foldl' _ v []      = v
foldl' f v (x:xs) = let z = f v x
                    in z `seq` foldl' f z xs
```

This version solves the problem with addition



Case study: foldl'




Figure 2:

[Faculty of Science
Information and Computing
Sciences]



Strict application

Most of the times we use `seq` to force an argument to a function, that is, *strict application*

```
($!) :: (a -> b) -> a -> b  
f $! x = x `seq` f x
```

Because of sharing, `x` is evaluated only once

```
foldl' _ v [] = v  
foldl' f v (x:xs) = ((foldl' f) $! (f v x)) xs
```



More (tricky) questions

What is the result of these expressions?

1. `(\x -> 0) $! undefined`
2. `seq (undefined, undefined) 0`
3. `snd $! (undefined, undefined)`
4. `(\x -> 0) $! (\x -> undefined)`
5. `undefined $! undefined`
6. `length $! map undefined [1,2]`
7. `seq (undefined + undefined) 0`
8. `seq (foldr undefined undefined) 0`
9. `seq (1 : undefined) 0`



More (tricky) questions

What is the result of these expressions?

1. `(\x -> 0) $! undefined = undefined`
2. `seq (undefined, undefined) 0 = 0`
3. `snd $! (undefined, undefined) = undefined`
4. `(\x -> 0) $! (\x -> undefined) = 0`
5. `undefined $! undefined = undefined`
6. `length $! map undefined [1,2] = 2`
7. `seq (undefined + undefined) 0 = undefined`
8. `seq (foldr undefined undefined) 0 = 0`
9. `seq (1 : undefined) 0 = 0`

seq only evaluates up to WHNF



Case study: Fibonacci numbers

```
fib 0 = 0
```


```
fib 1 = 1
```

```
fib n = fib (n-1) + fib (n-2)
```


What happens when we ask for `fib 5`?



Case study: Fibonacci numbers



Case study: Fibonacci numbers



Local memoization

Idea: remember the result for function calls

- ▶ We build a list of partial results
- ▶ Sharing takes care of evaluating only once

```
memo_fib n = map fib [0 .. ] !! n
  where fib 0 = 0
        fib 1 = 1
        fib n = memo_fib (n-1) + memo_fib (n-2)
```

You can get even faster by using a better data structure

- ▶ For example, IntMap from containers



Summary

- ▶ Laziness = evaluate only as much as needed
 - ▶ As opposed to the more common *eager* evaluation
- ▶ Evaluation is guided by pattern matching
 - ▶ We need WHNF to choose a branch
 - ▶ Some arguments may not even be evaluated
- ▶ Laziness is tricky when it fails
 - ▶ Too many thunks lead to a space leak
 - ▶ `seq` is used to *force* evaluation

